

Etymology-Driven Register Transfer in Romance L1 English: An Annotated Lexical and Corpus Resource

Darragh Kerins* Noumidia Abadour† Josu Barrutia‡ Mikel Molina§

University of the Basque Country (UPV/EHU) LAP10-LR, 2025–2026

*drussell1001@ikasle.ehu.eus †nabadour001@ikasle.ehu.eus

‡jbarrutia006@gmail.com §molinamikel10404@gmail.com

Abstract

English vocabulary is divided into Germanic and Latinate strata that carry distinct register signals. Romance L1 learners systematically favour the Latinate stratum in contexts where native speakers prefer Germanic alternatives, because Latinate words overlap with cognates in their first languages. This paper presents a pilot annotated resource comprising: (1) a 20-pair seed lexicon of Germanic–Latinate synonym pairs selected for register difference and Romance cognate overlap; (2) 38 manually annotated sentences drawn from EFCAMDAT, pre-labelled using Mistral 7B Instruct as a filtering oracle. Analysis of the full candidate pool ($N = 4,874$) and the annotated subset reveals that mismatch is lexically uneven, decreases consistently with CEFR proficiency level, and falls into two distinct error types: pure register mismatch and Romance meaning transfer via false cognates. Inter-annotator agreement on register markedness is high ($\kappa \geq 0.73$), while agreement on substitution naturalness is lower, reflecting the subjective nature of that judgement. The lexicon and annotated dataset are released as a reusable resource for learner language research and vocabulary instruction.

1 Introduction

The English lexicon contains two historically distinct strata. The *Germanic layer* entered Old English before the Norman Conquest of 1066 and constitutes the core of everyday, conversational usage: words such as *use*, *start*, *buy*, *help*, and *get*. The *Latinate layer* arrived via Anglo-Norman French after 1066 and became associated with formal, academic, and administrative registers: *utilise*, *commence*, *purchase*, *assist*, *obtain*. English is unusual cross-linguistically in maintaining this persistent register split within a single synonym set; most languages do not encode formality through etymological origin in the same systematic way (Leisi, 1985; Quirk et al., 1985).

For speakers of Romance languages, particularly Spanish and French, the Latinate stratum is immediately transparent. Words like *utilise*, *commence*, and *assist* overlap directly with their French equivalents (*utiliser*, *commencer*, *assister*) and Spanish equivalents (*utilizar*, *comenzar*, *asistir*). The corresponding Germanic words (*use*, *start*, *help*) have no Romance cognates, making them psycholinguistically less accessible to these learners.

The result is a **systematic register mismatch**: Romance L1 learners select Latinate vocabulary in neutral and informal contexts where native speakers prefer Germanic alternatives. The pattern is not random; it reflects a predictable interaction between L1 vocabulary structure and the English etymological register split. Despite being familiar to experienced teachers of English to Romance-speaking populations (Swan, 2005), the phenomenon has received limited attention as a structured corpus linguistics problem.

This paper makes the following contributions:

1. A **seed lexicon** of 20 Germanic–Latinate synonym pairs, selected against explicit etymological and register criteria.
2. A **pilot annotated corpus** of 38 learner sentences from EFCAMDAT, with four annotation dimensions per sentence.
3. An **LLM pre-annotation pipeline** using Mistral 7B Instruct as a filtering oracle, with a systematic profile of model behaviour before filtering.
4. A **qualitative error taxonomy** distinguishing pure register mismatch from Romance meaning transfer.

2 Background and Related Work

2.1 The English Etymological Register Split

Leisi (1985) documents the post-Conquest stratification of English vocabulary, noting that pairs such as *ask/inquire*, *begin/commence*, and *help/assist*

Context	Preferred form	Example
Informal/everyday	Germanic	<i>use</i> , not <i>utilise</i>
Neutral/conversational	Germanic	<i>start</i> , not <i>commence</i>
Formal/academic	Latinate	<i>acquire</i> , not <i>get</i>

Table 1: English register preferences by context.

are near-synonyms that differ primarily in stylistic weight rather than denotation. Quirk et al. (1985) further distinguishes a continuum from formal to informal register within such pairs, with the Germanic word almost always occupying the informal pole.

Table 1 illustrates the practical effect. In English, the choice between *use* and *utilise* signals register; in French and Spanish, the cognate forms (*utiliser/utilizar*) are neutral everyday vocabulary. Learners therefore approach English with a frequency-based bias that conflicts with English register norms.

2.2 Cross-Linguistic Similarity and Item Learning

Ringbom (1987) introduced the notion of *item learning* via cross-linguistic similarity: learners favour L2 forms that share surface properties with L1 words, acquiring them faster and using them more frequently, regardless of register or frequency in the target language. Ringbom (2007) refines this, noting that similarity facilitates comprehension more than production. In production, learners frequently transfer high-frequency L1 items to L2 contexts where they are stylistically inappropriate. For Spanish and French learners, Latinate English forms are cognatically accessible, and learners may be unaware of the register signal those forms carry in English.

2.3 Learner Corpora and Register

Granger (2015) surveys learner corpus research and notes a persistent tendency for advanced L2 writers to produce text that is lexically more formal than native speaker norms. Cobb and Horst (2007) links this to vocabulary knowledge: learners who know fewer high-frequency words compensate by relying on low-frequency Latinate forms acquired through reading. Crucially, this pattern means a learner can score highly on standard lexical sophistication measures while still producing systematic register errors: high-frequency Latinate forms such as *demonstrate* and *obtain* are not low-frequency in absolute terms, but they are over-formal for the conversational contexts in which learners use them.

Germanic	Latinate
begin	initiate
start	commence
buy	purchase
help	assist
need	require
ask	request
try	attempt
give	provide
show	demonstrate
end	terminate
use	utilise
get	obtain
tell	inform
see	observe
build	construct
keep	maintain
think	consider
put	place
find	locate
take	receive

Table 2: The 20 seed lexicon pairs, spanning transactions, communication, causation, and perception.

EFCAMDAT (Geertzen et al., 2014; Huang et al., 2017) provides the largest open corpus of learner writing with CEFR metadata, making it well-suited to studying proficiency effects on lexical register.

3 Resource Construction

3.1 Seed Lexicon

The seed lexicon consists of 20 Germanic–Latinate synonym pairs selected using four criteria:

1. The Latinate word has a direct cognate in both Spanish and French.
2. The Germanic word has no close Romance equivalent.
3. The pair differs primarily in register, not denotation.
4. Etymology was verified via the Oxford English Dictionary.

Selected pairs are shown in Table 2. Pairs span multiple semantic fields (transactions, communication, causation, perception) to separate lexical-field effects from register effects.

3.2 Corpus: EFCAMDAT

Source sentences were drawn from EFCAMDAT (Geertzen et al., 2014; Huang et al., 2017), a corpus of English writing produced by learners worldwide on the EF platform, containing approximately 70 million words. We targeted French and Spanish L1 writers across CEFR levels A1–C1.

3.3 Sampling and Filtering Pipeline

The pipeline proceeded as follows:

1. Filter to French and Spanish L1, CEFR levels A1–C1.
2. Sample up to 10,000 essays per L1×CEFR cell (random seed 42).
3. Apply the 20-lemma lexicon filter, matching base forms and inflected variants (-s, -ed, -ing, -ise/-ize).
4. Pass each candidate sentence to Mistral 7B Instruct (Jiang et al., 2023) (local deployment via Ollama) for pre-annotation of four fields: `register_label`, `register_markedness`, `simpler_alternative`, and `substitution_natural`, plus a confidence score from 1 (low) to 3 (high).
5. Retain candidates with confidence ≥ 2 and `register_markedness` = 1; cap at 2 rows per lemma to ensure lexical balance.

This produced a final human-annotation set of **38 sentences**.

3.4 Annotation Dimensions

Each sentence was annotated on four dimensions:

- **Context label:** one of {informal, neutral essay, academic, news, ambiguous}.
- **Register markedness:** binary; 0 = appropriate for context, 1 = overly formal.
- **Simpler alternative:** the suggested Germanic replacement.
- **Substitution naturalness:** binary; 0 = unnatural, 1 = natural.

3.5 Annotation Team

The four annotators represent complementary linguistic profiles (Table 3): a native English speaker (Darragh), a speaker with French and Amazigh background (Noumidia), and two Spanish L1 speakers (Josu, Mikel), providing both a native English baseline and genuine target-learner perspectives.

4 Results

4.1 LLM Pre-annotation Profile

Before filtering, the full candidate pool contained **4,874 sentences**. Mistral 7B flagged 51.7% as over-formal, indicating that the model does not trivially accept everything as mismatched. The majority of candidates (63%) came from neutral essay contexts,

showing that register mismatch is not confined to informal writing.

4.2 CEFR Distribution

Figure 1 shows over-formality rates by CEFR level. Register-marked candidates fell from 61% at A1 to 35% at C1, a consistent monotonic decrease. Lower-proficiency learners produce the most detectable mismatch, and the pattern holds across both L1 groups.

4.3 Lemma-Level Mismatch Rates

Table 4 shows the LLM-judged over-formal ratio per lemma, i.e. the proportion of sentences where the Latinate word was judged register-marked. Mismatch rates vary considerably across lexical items, directly addressing RQ2.

Demonstrate (0.90) and *obtain* (0.67) are flagged in nearly all contexts. *Place* has a high raw frequency (3,205 hits) but a lower mismatch ratio (0.38), because it appears in collocations such as *take place* where it is not register-marked. The 2-per-lemma cap prevents *place* from dominating the annotated set. *Commence* produced zero corpus hits, confirming it lies above the productive vocabulary threshold for A1–C1 learners.

4.4 Qualitative Error Taxonomy

Manual review of the 38 annotated sentences reveals two structurally distinct error types:

Type 1: Register mismatch. The Latinate word is semantically correct but stylistically too heavy for the context.

**I utilise my phone to listen to music. → I use my phone...*

**We purchased some food at the market. → We bought...*

**I require a new notebook. → I need a new notebook.*

Type 2: False cognate/meaning transfer. The learner has borrowed not just the form but also a Romance meaning that does not carry over into English. The error is both a register mismatch and a semantic one.

**I assisted at the match. (Fr. assister = attend)*

**My teacher initiated me to painting. (Fr. initier = introduce to)*

**We terminated the project. (Sp. terminar = finish; also over-formal)*

Type 1 calls for register awareness training; Type 2 calls for explicit false-cognate instruction.

Annotator	L1	Role
Darragh Kerins	Irish/English	Native English baseline
Noumidia Abadour	Amazigh/Arabic	Romance competence (French, Spanish)
Josu Barrutia	Spanish	Target learner perspective
Mikel Molina	Spanish	Target learner perspective

Table 3: Annotation team linguistic profiles.




Figure 1: Proportion of over-formal candidates by CEFR level (full pre-filter pool, $N = 4,874$).

Lemma	Over-formal ratio	Corpus hits
<i>demonstrate</i>	0.90	21
<i>obtain</i>	0.67	30
<i>inform</i>	0.55	58
<i>assist</i>	0.55	44
<i>purchase</i>	0.52	25
<i>receive</i>	0.47	189
<i>observe</i>	0.45	22
<i>consider</i>	0.45	164
<i>construct</i>	0.42	19
<i>require</i>	0.40	97
<i>request</i>	0.39	51
<i>place</i>	0.38	3205
<i>maintain</i>	0.33	73
<i>provide</i>	0.32	148

Table 4: Over-formal ratios and raw corpus hit counts per lemma (full pre-filter pool).

4.5 Inter-Annotator Agreement

Figure 2 shows human–human pairwise agreement across the three annotation dimensions. Agreement on register markedness averaged 94%, and all pairwise Cohen’s κ values met or exceeded 0.70 (Artstein and Poesio, 2008), as shown in Figure 3. Agreement on substitution naturalness was lower at 74%, which is expected given the more subjective nature of that judgement.

Human–AI register agreement ranged from 71–88%, partly a selection effect since the dataset was

pre-filtered by the same model on that question. Human–AI naturalness agreement was only 35–47%, showing that Mistral 7B does not reliably judge whether a Germanic substitution sounds natural.

Figure 4 plots each annotator’s register strictness against naturalness acceptance. Darragh applied the most lenient register threshold and was most accepting of Germanic substitutions. Josu and Mikel were stricter on register but more conservative about naturalness, consistent with lower exposure to colloquial Germanic vocabulary. Noumidia’s profile is the most balanced across both dimensions.

5 Discussion

5.1 Lexical Asymmetry in Register Mismatch

The mismatch rates in Table 4 show that Latinate verbs are not equally register-marked. *Demonstrate* and *obtain* are almost always over-formal in learner essays, while *provide* and *maintain* show weaker, more context-dependent effects. Pedagogical intervention should prioritise high-mismatch pairs rather than treating the Germanic/Latinate distinction as a blanket rule.




Figure 2: Human–human pairwise agreement rates across the three annotation dimensions (first 10 shared examples).




Figure 3: Pairwise Cohen’s κ on register markedness. All values meet the 0.70 reliability threshold.

5.2 Proficiency and Register Calibration

The monotonic decrease in Figure 1 is consistent with Ringbom (2007)’s account of learners shifting from item learning towards system learning as proficiency grows. The rate at C1 (35%) shows, however, that register calibration is not complete at near-advanced levels, in line with Granger (2015)’s observation that over-formality persists across proficiency bands.

5.3 The LLM as Filtering Oracle

The gap between human–AI register agreement (82%) and naturalness agreement (43%) is itself informative. Register markedness is a relatively coarse binary judgement: a word is or is not over-formal for a given context. Substitution naturalness is a more fine-grained pragmatic judgement that

requires sensitivity to colloquial usage, idiomatic preference, and spoken register. The results suggest that Mistral 7B has a working model of formal/informal contrast but lacks the native-speaker intuition needed to judge whether a simpler word *sounds right* in a given sentence. This distinction between formal register detection and pragmatic naturalness is relevant beyond this pipeline: it suggests that LLMs may be unreliable evaluators of substitution quality in any task requiring colloquial intuition. Future pipelines should use the LLM only for the register labelling step and treat naturalness as a human-only annotation task from the outset.




Figure 4: Annotator profiles: register strictness (x-axis) against naturalness acceptance (y-axis).

6 Conclusion

This paper presents a pilot lexical resource for studying etymology-driven register transfer in Romance L1 English. The core finding is that register mismatch is not uniform: it is concentrated in a subset of Latinate verbs (*demonstrate*, *obtain*, *inform*) that are nearly always over-formal in learner essays, while others (*provide*, *maintain*) show context-dependent effects. Mismatch decreases monotonically from A1 to C1 but remains detectable at near-advanced levels, suggesting that register calibration is not a by-product of general proficiency growth alone. The error taxonomy identifies two distinct phenomena, pure register mismatch and false-cognate meaning transfer, which require different pedagogical responses.

The pipeline demonstrates that a small-instruction LLM can serve as a useful pre-filter for register markedness, reducing the annotation workload without introducing systematic bias. However, the 43% human–AI agreement on substitution naturalness shows that pragmatic naturalness judgements require human annotation. The resource, though small, provides a reusable foundation: the seed lexicon and annotation schema can be extended to larger samples, additional L1 groups, and spoken corpora.

Limitations

Dataset size. The annotated set contains 38 sentences. Findings are indicative and should not be over-generalised. L1 comparisons between French

and Spanish (RQ1) require a substantially larger balanced sample.

CEFR via placement scores. Learner levels are derived from EFCAMDAT placement tests rather than independently verified proficiency assessments, and may not align consistently with standardised CEFR benchmarks.

Essay genre. EFCAMDAT texts are learner essays, already moderately formal. This likely suppresses informal register choices and may underestimate mismatch in spoken or informal written contexts.

LLM naturalness judgements. Human–AI agreement on naturalness was only ~43%. Mistral 7B labels on this dimension should not replace human annotation without further prompt refinement and validation.

Lexicon coverage. The 20-pair lexicon covers frequent verbs only. The register split extends to adjectives, nouns, and adverbs (*glad/content*, *kingly/regal*), which are not addressed here.

Data Statement

The annotated dataset and seed lexicon are available on request. EFCAMDAT source data is available at <https://ef-lab.mml1.cam.ac.uk/EFCAMDAT.html>. All sentences are learner-produced texts from the public corpus; no personally identifiable information is included. The annotation was conducted in 2025–2026 at UPV/EHU under the Building Language Resources module (LAP10-LR).

Acknowledgements

We thank Izaskun Aldezabal Roteta for guidance throughout the LAP10-LR module at UPV/EHU and for feedback on earlier versions of this work.

References

- Ron Artstein and Massimo Poesio. Inter-coder agreement for computational linguistics. *Computational Linguistics*, 34(4):555–596, 2008. doi: 10.1162/coli.07-034-R2.
- Tom Cobb and Marlise Horst. Understanding Web-based concordancing for CALL. *CALICO Journal*, 24(2):285–308, 2007.
- Jeroen Geertzen, Theodora Alexopoulou, and Anna Korhonen. Automatic linguistic annotation of large scale L2 databases: The EF-Cambridge open language database (EFCamDat). In Ron T. Millar, Kate I. Martin, Casey M. Eddington, Allison Henery, Nausica M. Miguel, and Amanda

Tseng, editors, *Selected Proceedings of the 2012 Second Language Research Forum*, pages 240–254, Somerville, MA, 2014. Cascadilla Proceedings Project.

Sylviane Granger. Learner corpora: Past, present and future. In Sylviane Granger, Gaëtanelle Gilquin, and Fanny Meunier, editors, *The Cambridge Handbook of Learner Corpus Research*, pages 1–5. Cambridge University Press, Cambridge, 2015.

Yan Huang, Jeroen Geertzen, Russell Baker, Anna Korhonen, and Theodora Alexopoulou. The EF Cambridge open language database (EFCAMDAT): Information for users. Technical report, EF Education First Research Lab, University of Cambridge, 2017.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Léo Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. Mistral 7B. *arXiv preprint arXiv:2310.06825*, 2023.

Ernst Leisi. *Der Wortinhalt: Seine Struktur im Deutschen und Englischen*. Quelle & Meyer, Heidelberg, 6 edition, 1985.

Randolph Quirk, Sidney Greenbaum, Geoffrey Leech, and Jan Svartvik. *A Comprehensive Grammar of the English Language*. Longman, London, 1985.

Håkan Ringbom. *The Role of the First Language in Foreign Language Learning*. Multilingual Matters, Clevedon, 1987.

Håkan Ringbom. *Cross-Linguistic Similarity in Foreign Language Learning*. Multilingual Matters, Clevedon, 2007.

Michael Swan. *Practical English Usage*. Oxford University Press, Oxford, 3 edition, 2005.